

The same learner, the same splits — only the decomposition protocol differs

decompose-then-split reports 45.9 % lower error than is attainable causally: the gain is an artefact of the protocol, not of the model

