


Decompose-then-split reports 46–51 % lower error on both benchmarks, with the protocol as the only difference

“matched”: the same CEEMDAN (20 trials, $K = 8$), the same feature layout and the same ridge, fitted over the whole series or recomputed inside each input window.

“original”: the contrast as first run, which also swaps the decomposition family and the fitting scheme (GEFCom2014 45.9 %, PJM 52.4 %) and therefore bounds the effect rather than isolating it.

GEFCom2014 · protocol accounts for 46.4 %



PJM · protocol accounts for 50.9 %

