

Holding everything but the protocol fixed, decompose-then-split reports 46.4 % lower error

“matched”: the same CEEMDAN (20 trials, K = 8), the same feature layout and the same ridge, fitted over the whole series or recomputed inside each window — the protocol is the only difference, and it accounts for 46.4 %. “original”: the contrast as first run, which also swaps the decomposition family and the fitting scheme and gives 45.9 %; it bounds the effect rather than isolating it.

